



## Capstone Project

Edge computing device programming for AI projects

### Lecture 4

Dr. Wilton Fok &

Carol Chen

# Recap: Training classification model

	Operation	Results
Dataset preparation	(Preparing the catdog classification dataset)	Get a prepared dataset
Train model	<i>cd ~/jetson-inference/python/training/classification python3 train.py --model-dir=models/cat_dog data/cat_dog -- batch-size=1 --workers=1 --epochs=1</i>	Obtaining a .pth model
Convert model	<i>python3 onnx_export.py --input=model_best.pth.tar --model-dir=models/cat_dog</i>	Obtaining a .onnx model
Run model	<i>/home/nvidia/jetson-inference/build/aarch64/bin/imagenet -- model= models/resnet18.onnx --input_blob=input_0 -- output_blob=output_0 --labels=data/cat_dog/labels.txt data/cat_dog/test/cat</i>	Visualize the output of the model

# GRAMMAR and TRANSLATION of terminal command

Blue: predicate (verb)

Green: object (noun)


Purple: adverbial // arguments

	Operation	Translation
Entering target folder	<code>cd ~/jetson-inference/python/training/classification</code>	Told the system to <i>Change the directory</i> into <code>~/jetson-inference/python/training/classification</code>
Train model	<code>python3 train.py --model-dir=models/cat_dog data/cat_dog --batch-size=1 --workers=1 --epochs=1</code>	Use the <i>software python3</i> to run the <i>file "train.py"</i> with the following arguments " <code>--model-dir=models/cat_dog data/cat_dog --batch-size=1 --workers=1 --epochs=1</code> "
Convert model	<code>python3 onnx_export.py --input=model_best.pth.tar --model-dir=models/cat_dog</code>	Use the <i>software python3</i> to run the file <i>onnx_export.py</i> using the model <code>models/cat_dog</code>
Run model	<code>/home/nvidia/jetson-inference/build/aarch64/bin/imagenet --model=models/resnet18.onnx --input_blob=input_0 --output_blob=output_0 --labels=data/cat_dog/labels.txt data/cat_dog/test/cat save/cat</code>	Directly Execute the file <code>"/home/nvidia/jetson-inference/build/aarch64/bin/imagenet"</code> with the following arguments " <code>--model=models/resnet18.onnx --input_blob=input_0 --output_blob=output_0 --labels data/cat_dog/labels.txt data/cat_dog/test/cat save/cat</code> "

For each elements, we should add a space to separate them; While finish a sentence, using Enter for ending.

# What's the arguments?

Argument means the parameters you use while executing something.




# What's the arguments?

## Command in terminal

```
python3 train.py --model-dir=models/cat_dog data/cat_dog --batch-size=1 --workers=1 --epochs=1
```

```
/home/nvidia/jetson-inference/build/aarch64/bin/imagenet --model= models/resnet18.onnx --input_blob=input_0 --  
output_blob=output_0 --labels=data/cat_dog/labels.txt data/cat_dog/test/cat save/cat
```

This library allows us to pass parameters into the program directly from the command line. We can use python file.py to run python files. The function of argparse is to parse, save and use other parameters passed in from the command line.



## .py file

```
parser.add_argument("input_URI", type=str, default="", nargs='?', help="URI of the input stream")  
parser.add_argument("output_URI", type=str, default="", nargs='?', help="URI of the output stream")  
parser.add_argument("--network", type=str, default="resnet18-body", help="pre-trained model to load (see below for options)")  
parser.add_argument("--overlay", type=str, default="links,keypoints", help="pose overlay flags (e.g. --overlay=links,keypoi")  
parser.add_argument("--threshold", type=float, default=0.15, help="minimum detection threshold to use")
```

```
# parser.add_argument("--keypoint1", type=str, default="", help="when calculating the angle, choose the first keypoint")  
# parser.add_argument("--keypoint2", type=str, default="", help="when calculating the angle, choose the second keypoint")  
# parser.add_argument("--keypoint3", type=str, default="", help="when calculating the angle, choose the third keypoint")
```

# What's the arguments?

- Argument means the parameters you use while executing something.

Command	Translation
<code>python3 train.py --model-dir=models/cat_dog data/cat_dog --batch-size=1 --workers=1 --epochs=1</code>	<i>we are telling the system that for this training, we need to “store the model in models/cat_dog, using the data from data/cat_dog. The training uses 1 batchsize and 1 workers. We train the model for 1 epochs.”</i>
<code>/home/nvidia/jetson-inference/build/aarch64/bin/imagenet --model=models/resnet18.onnx --input_blob=input_0 --output_blob=output_0 --labels=data/cat_dog/labels.txt data/cat_dog/test/cat save/cat</code>	<i>we are telling the system that while running this file, “I am using the model from models/resnet18.onnx, and its input_blob and output_blob are input_0 and output_0 respectively. Also, I want to processing the images in the data/cat_dog/test/cat folder, and the corresponding label is data/cat_dog/labels.txt data/cat_dog/test/cat. Please save the processed images into save/cat”</i>

Grammar: For each argument, we use double dash -- **to specify**, = **to assign**, and **space for termination**.

For example, we want to set model-dir to be “models/self”, batch-size to be 1, then the command will be:

**--model-dir=models/self --batch-size=1**

*Remember a space here* 


# How to achieve AI training?

training      *python3 train.py --model-dir=models/cat\_dog data/cat\_dog --batch-size=1 --workers=1 --epochs=1*


inference      */home/nvidia/jetson-inference/build/aarch64/bin/imagenet --model= models/resnet18.onnx --input\_blob=input\_0 --output\_blob=output\_0 --labels=data/cat\_dog/labels.txt data/cat\_dog/test/cat save/cat*



# Pose estimation







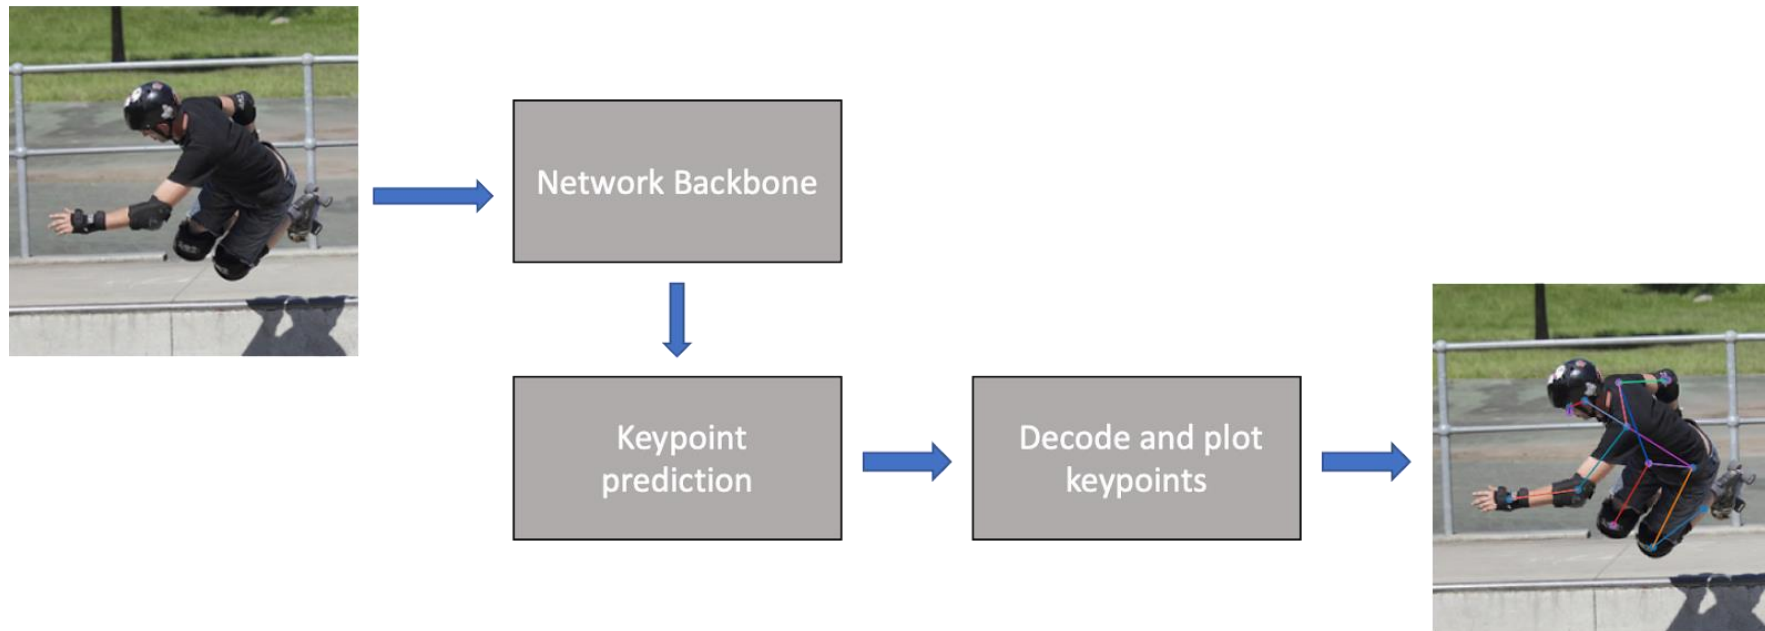
## Pose estimation- Introduction

Human Pose Estimation identifies the poses of human body parts and joints in images or videos.

# Pose estimation- Introduction

---

When an image or video is given to the pose estimator model as input, it identifies the coordinates of the detected body parts like wrist, shoulder, knees, eyes, arms and so on, then joints as output and a confidence score showing precision of the estimations.




The architecture of Human Pose Estimation

# Pose estimation-Application


Human pose estimation is widely used in many applications, here are a few common ones:



Human-computer interaction, by analyzing human posture and movements, enables non-contact human-computer interaction.



Motion analysis can analysis the human body's pose and movements to analyze the body's movement state and trajectory, and do the motion correction



Medical Rehabilitation helps Patients with Stroke Recovery by estimating their movements




# Different kinds of pose estimation

- 2D pose estimation

2D pose estimation uses a single camera or image to predict the location of keypoints.

There are two representations of keypoint locations: **coordinate** and **heatmap**. The coordinate representation only needs to regress the location of each keypoint, while the heatmap representation requires the network to predict the probability that each pixel point belongs to a certain type of keypoint.



Coordinate representation



(a)



(b)




(c)

Heatmap representation

- 3D pose estimation

3D pose estimation uses multiple cameras or image sequences to infer the 3D pose of the human body. Such methods usually require camera calibration and correlation of multiple viewpoints to recover the 3D pose of the human body. There are two main forms of output of 3D HPE prediction:

1) 3D skeleton, which is similar to 2D, except that the skeleton information is presented in a three-dimensional coordinate system




2) 3D shape or mesh, which uses either a parametric human body model or a triangulated mesh to present the predicted results, and which is a more dense skeleton information



# Some popular approaches for pose estimation

- **SMPL: A Skinned Multi-Person Linear Model (2015)**

A standard 3D human model, the main parameters are shape blend shapes parameters to control the body shape and pose blend shapes parameters to control the body posture.(video link: [https://www.youtube.com/watch?v=kuBUyHeV5U&t=287s&ab\\_channel=MichaelBlack](https://www.youtube.com/watch?v=kuBUyHeV5U&t=287s&ab_channel=MichaelBlack))



# Some popular approaches for pose estimation

- **Object-Occluded Human Shape and Pose Estimation from a Single Color Image (2020)**

This research is using the SMPL model as an output representation to extract the 3d mesh of the SMPL model from a single map and also solves the problem of occlusion (video link: [https://www.youtube.com/watch?v=8udm2OB0A-U&t=189s&ab\\_channel=YangangWang](https://www.youtube.com/watch?v=8udm2OB0A-U&t=189s&ab_channel=YangangWang))





# Some popular approaches for pose estimation

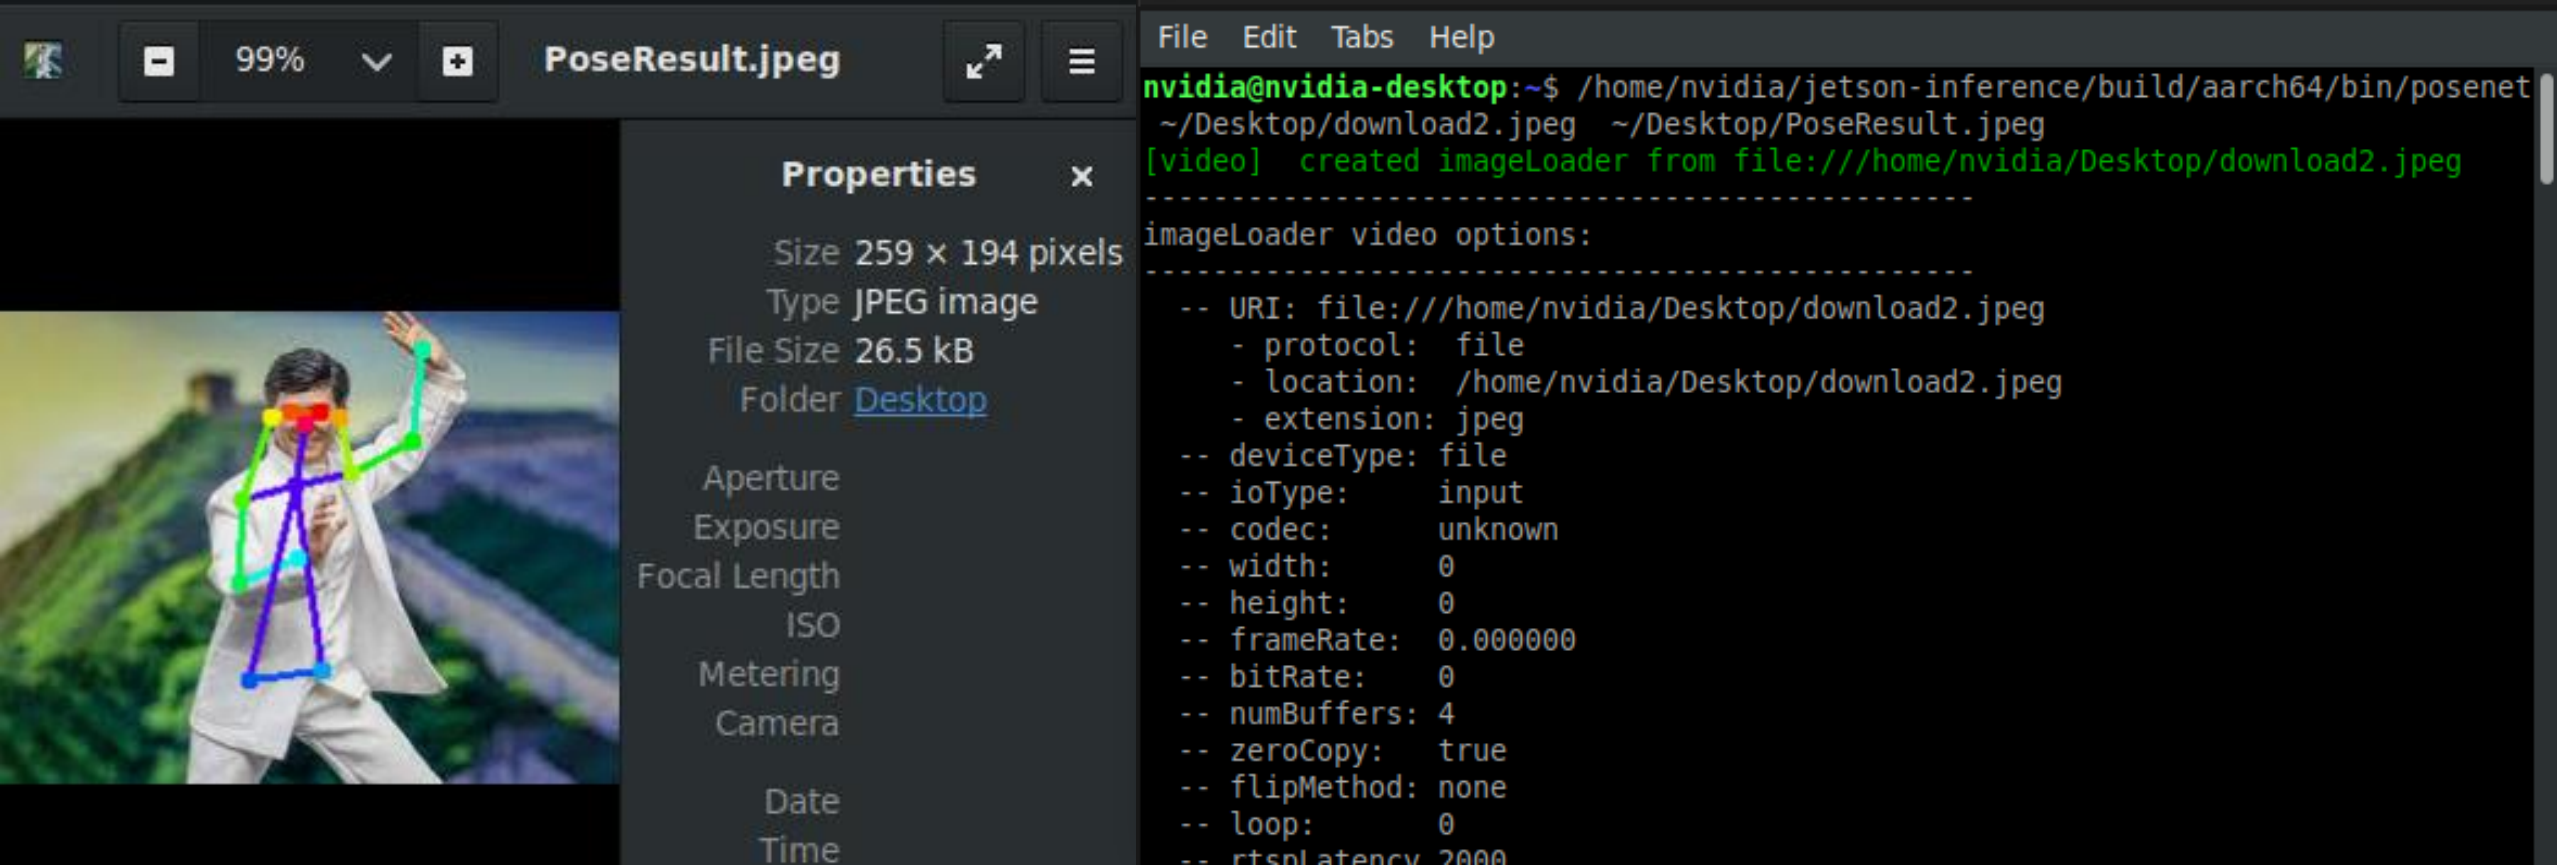
- **RMPE: Regional Multi-Person Pose Estimation (2017) AlphaPose**

RMPE can facilitate pose estimation in the presence of inaccurate human bounding boxes. The approach follows the two-step framework, first detects human bounding boxes and then estimating the pose within each box independently. (video link: [https://www.youtube.com/watch?v=Z2WPd59pRi8&ab\\_channel=haoshufang](https://www.youtube.com/watch?v=Z2WPd59pRi8&ab_channel=haoshufang))




# Pose estimation --- Image

- `$ /home/nvidia/jetson-inference/build/aarch64/bin/posenet YourImagePath ResultSavePath`



# Pose estimation --- Video

- `$ /home/nvidia/jetson-inference/build/aarch64/bin/posenet /dev/video0`  
OR
- `$ cd /home/nvidia/jetson-inference/build/aarch64/bin`
- `$ posenet /dev/video0`



# How to get the coordinate of keypoints?

Use the pretrained model Pose-ResNet18-Body

Input an image, we can get the object bbox, keypoints and links as output.

```
# load the pose estimation model
```

```
net = jetson.inference.poseNet(opt.network, sys.argv, opt.threshold)
```

```
# perform pose estimation (with overlay)
```

```
poses = net.Process(img, overlay=opt.overlay)
```

```
for pose in poses:
```

```
    print(pose.Keypoints) #print keypoints id, x, y. (x,y) is the coordinate
```





# Use PoseNet to get the keypoints and an angle

---

**If we want to detect the bending angle of the right arm of this lady, how should we implement it in code?**

We have 18 keypoints: ["nose", "left\_eye", "right\_eye", "left\_ear", "right\_ear", "left\_shoulder", "right\_shoulder", "left\_elbow", "right\_elbow", "left\_wrist", "right\_wrist", "left\_hip", "right\_hip", "left\_knee", "right\_knee", "left\_ankle", "right\_ankle", "neck"]

# Use PoseNet to get the keypoints and an angle

Download [new-posenet.py](#) [test\\_image](#) and put them in the path:  
/home/nvidia/jetson-inference/python/examples

Open a terminal and use these commands:

```
$cd /home/nvidia/jetson-  
inference/python/examples
```

```
$python3 new-posenet.py test.jpg  
save.jpg
```

In this case, we calculate the angle  
right\_shoulder-right\_elbow-right\_wrist

```
nvidia@nvidia-desktop: ~/jetson-inference/python/examples
-- x: 440.034
-- y: 774.678
, <poseNet.ObjectPose.Keypoint object>
-- ID: 14 (right_knee)
-- x: 336.785
-- y: 781.461
, <poseNet.ObjectPose.Keypoint object>
-- ID: 15 (left_ankle)
-- x: 401.22
-- y: 929.378
, <poseNet.ObjectPose.Keypoint object>
-- ID: 16 (right_ankle)
-- x: 348.329
-- y: 967.715
, <poseNet.ObjectPose.Keypoint object>
-- ID: 17 (neck)
-- x: 368.814
-- y: 359.717
]
Links [(13, 15), (11, 13), (14, 16), (12, 14), (11, 12), (5, 7), (6, 8), (7, 9),
(8, 10), (1, 2), (0, 1), (0, 2), (1, 3), (2, 4), (3, 5), (4, 6), (0, 17), (5, 1
7), (6, 17), (11, 17), (12, 17)]
Angle between right shoulder-right elbow-right wrist is: 88.41
```

# How can we utilise the keypoint to do the calculation?

```
def calculate_angle(id1, id2, id3):  
    point1_x = np.array(pose.Keypoints[id1].x)  
    point1_y = np.array(pose.Keypoints[id1].y)  
    point2_x = np.array(pose.Keypoints[id2].x)  
    point2_y = np.array(pose.Keypoints[id2].y)  
    point3_x = np.array(pose.Keypoints[id3].x)  
    point3_y = np.array(pose.Keypoints[id3].y)  
    v1 = (point1_x - point2_x, point1_y - point2_y)  
    v2 = (point3_x - point2_x, point3_y - point2_y)  
    cosine_angle = np.dot(v1, v2) / (np.linalg.norm(v1) * np.linalg.norm(v2))  
    angle = np.arccos(cosine_angle)  
    return np.degrees(angle)
```

Calculate the angle among 3 keypoints:

Get three keypoints' coordinates, and calculate the two vectors:

$$\theta = \cos^{-1} [ (a \cdot b) / (|a| |b|) ]$$

The python code is:

```
cosine_angle = np.dot(v1, v2) /  
    (np.linalg.norm(v1) * np.linalg.norm(v2))  
angle = np.arccos(cosine_angle)  
return np.degrees(angle)
```




# Object segmentation

---



# Object segmentation --- Video



Download [model](#) and put it in `/home/nvidia/jetson-inference/build/aarch64/bin/networks/`  
Then run the command in the terminal:

```
$ /home/nvidia/jetson-inference/build/aarch64/bin/segnet --network=fcn-resnet18-voc  
/dev/video0
```



# Assignment 4

- Detect the angle of the left arm
- Detect the angle of the right leg
- Detect the angle of the left leg
- Show the angle value on the image

Input image: [test\\_image](#)

Deadline: the end of the today's course (9:30pm)

Assessment in class by TA



# Action Recognition

---



# Action Recognition in Videos

Action Recognition in Videos is a task in computer vision where the goal is to **identify and categorize human actions performed in a video sequence**. The task involves analyzing the spatiotemporal dynamics of the actions and mapping them to a predefined set of action classes, such as running, jumping, or swimming.



push



pushup



ride  
bike



ride  
horse



run



shake  
hands



shoot  
ball



shoot  
bow



shoot  
gun



sit



situp



smile




smoke



somersault



# Demo of Action Recognition



[https://youtu.be/7\\_mcWCB76Ps?list=PL18uP4AzeyVuqH7dRNmyLGR1USnSgfliv](https://youtu.be/7_mcWCB76Ps?list=PL18uP4AzeyVuqH7dRNmyLGR1USnSgfliv)

# Run ActionNet

- Download the demo video and put it in the path:
- `$ cd Jeston-inference/python/examples`
- `$ python3 actionnet.py demo.mp4 result.mp4`



# Run ActionNet-Camera

- `$ cd Jeston-inference/python/examples`
- `$ python3 actionnet.py /dev/video0`

# Background Removal

---






## Background Removal

Background removal (aka background subtraction, or salient object detection) **generates a mask that segments the foreground from the background** of an image. It can be used to replace or blur backgrounds (similar to video conferencing applications), or it could aid in pre-processing for other vision DNN's like object detection/tracking or motion detection.

# Run BackgroundNet



- Download the demo image and put it in the path:
- `$ cd Jeston-inference/python/examples`
- `$ python3 backgroundnet.py images/bird_0.jpg images/test/bird_mask.png`  
`# remove the background (with alpha)`
- `$ python3 backgroundnet.py --replace=images/snow.jpg images/bird_0.jpg images/test/bird_replace.jpg` `# replace the background`


# Run BackgroundNet-Camera

- `$ cd Jeston-inference/python/examples`
- `$ python3 backgroundnet /dev/video0` # remove the background
- `$ python3 backgroundnet --replace=images/coral.jpg /dev/video0`  
# replace the background




# Monocular Depth Estimation

---



**Monocular Depth Estimation** is the task of estimating the **depth value (distance relative to the camera) of each pixel** given a single (monocular) RGB image. This challenging task is a key prerequisite for determining scene understanding for applications such as 3D scene reconstruction, autonomous driving, and AR.



[Real time Depth Anything V2 demo](#)



# Mono Depth on Images

- `$ cd Jeston-inference/python/examples`
- `$ python3 depthnet.py "images/room_*.jpg" images/test/depth_room_%i.jpg`



# Mono Depth from video

- `$ cd Jetson-inference/python/examples`
- `$ python3 depthnet.py /dev/video0`

